

Name: Jin En Tan

Student ID:31336574

FIT3152 Assignment 2

## Report

### Question 1:

```
r<=WAUS %>% count(warmerTomorrow) #ignoring those NA
  warmerTomorrow      n
1              0    887
2              1   1094
3             <NA>    19
```

Based on the above screenshot from R-studio console, we can see the number of days which is warmer than previous day indicated by 1 and number of days which is cooler than previous day indicated by 0. We will ignore the number of days with NA as only 19 rows are with NA so it will be less likely to affect the analysis.

```
r<=887/(887+1094)*100 #this is percentage of cooler than previous day
[1] 44.8
r<=1094/(887+1094)*100 #this is percentage of warmer than previous day
[1] 55.2
```

Based on the above screenshot, the proportion of days when it is warmer than the previous day is 55.2% whereas the proportion of days when it is cooler than the previous day is 44.8%. We can see that the percentage of days warmer than previous day is higher than that of cooler than previous day.

```
r<=summary(WAUS)
  Day      Month      Year      Location      MinTemp      MaxTemp      Rainfall      Evaporation      Sunshine
Min. : 1.0 Min. : 1.00 Min. :2008 Min. : 7.0 Min. : -2.4 Min. : 9.4 Min. : 0.0 Min. : 0 Min. : 0
1st Qu.: 8.0 1st Qu.: 4.00 1st Qu.:2011 1st Qu.:19.0 1st Qu.: 7.3 1st Qu.:17.1 1st Qu.: 0.0 1st Qu.: 2 1st Qu.: 5
Median :16.0 Median : 7.00 Median :2014 Median :23.0 Median :11.2 Median :21.9 Median : 0.0 Median : 4 Median : 8
Mean :15.7 Mean : 6.73 Mean :2014 Mean :23.6 Mean :11.4 Mean :22.6 Mean : 1.9 Mean : 5 Mean : 8
3rd Qu.:23.0 3rd Qu.:10.00 3rd Qu.:2017 3rd Qu.:31.0 3rd Qu.:15.5 3rd Qu.:27.0 3rd Qu.: 0.6 3rd Qu.: 7 3rd Qu.:11
Max. :31.0 Max. :12.00 Max. :2019 Max. :34.0 Max. :26.6 Max. :44.1 Max. :143.8 Max. :21 Max. :14
NA's :19 NA's :19 NA's :19 NA's :21 NA's :14 NA's :47 NA's :1025 NA's :1095
  WindGustDir WindGustSpeed WindDir9am WindDir3pm WindSpeed9am WindSpeed3pm Humidity9am Humidity3pm Pressure9am
SSE : 186 Min. : 9.0 N : 219 S : 182 Min. : 0.0 Min. : 0.0 Min. : 10.0 Min. : 4.0 Min. : 994
W : 185 1st Qu.: 30.0 SSE : 166 W : 180 1st Qu.: 9.0 1st Qu.:13.0 1st Qu.: 57.0 1st Qu.: 36.0 1st Qu.:1013
N : 168 Median : 39.0 S : 164 WSW : 159 Median :13.0 Median :19.0 Median : 71.0 Median : 53.0 Median :1018
S : 152 Mean : 40.7 NNE : 136 N : 151 Mean :15.1 Mean :19.7 Mean : 69.8 Mean : 51.8 Mean :1018
SW : 149 3rd Qu.: 50.0 SSW : 135 SW : 149 3rd Qu.:20.0 3rd Qu.:24.0 3rd Qu.: 84.0 3rd Qu.: 67.0 3rd Qu.:1023
(Other):1115 Max. :111.0 (Other):1081 (Other):1147 Max. :63.0 Max. :61.0 Max. :100.0 Max. :100.0 Max. :1037
NA's : 45 NA's :37 NA's : 99 NA's : 32 NA's :27 NA's :24 NA's :30 NA's :243
  Pressure3pm Cloud9am Cloud3pm Temp9am Temp3pm warmerTomorrow
Min. : 992 Min. :0 Min. :0 Min. : 0.9 Min. : 5.0 0 : 887
1st Qu.:1012 1st Qu.:1 1st Qu.:2 1st Qu.:11.7 1st Qu.:16.0 1 :1094
Median :1016 Median :6 Median :6 Median :15.6 Median :20.4 NA's : 19
Mean :1016 Mean :5 Mean :5 Mean :16.0 Mean :21.0
3rd Qu.:1021 3rd Qu.:7 3rd Qu.:7 3rd Qu.:20.1 3rd Qu.:25.2
Max. :1035 Max. :8 Max. :8 Max. :35.3 Max. :43.0
NA's :246 NA's :934 NA's :927 NA's :21 NA's :23
```

Based on the descriptions of the predictor variables, we can see that the noteworthy thing in the data is that for real-value attributes, there is a lot of NA's in the data. We can see that evaporation, sunshine, cloud9am and cloud3pm contains a lot of NA's which is almost or already exceeding half the data size of 2000. Hence, I consider to remove these 4 attributes from my analysis as there's too many NA's in the data.

### **Question 2:**

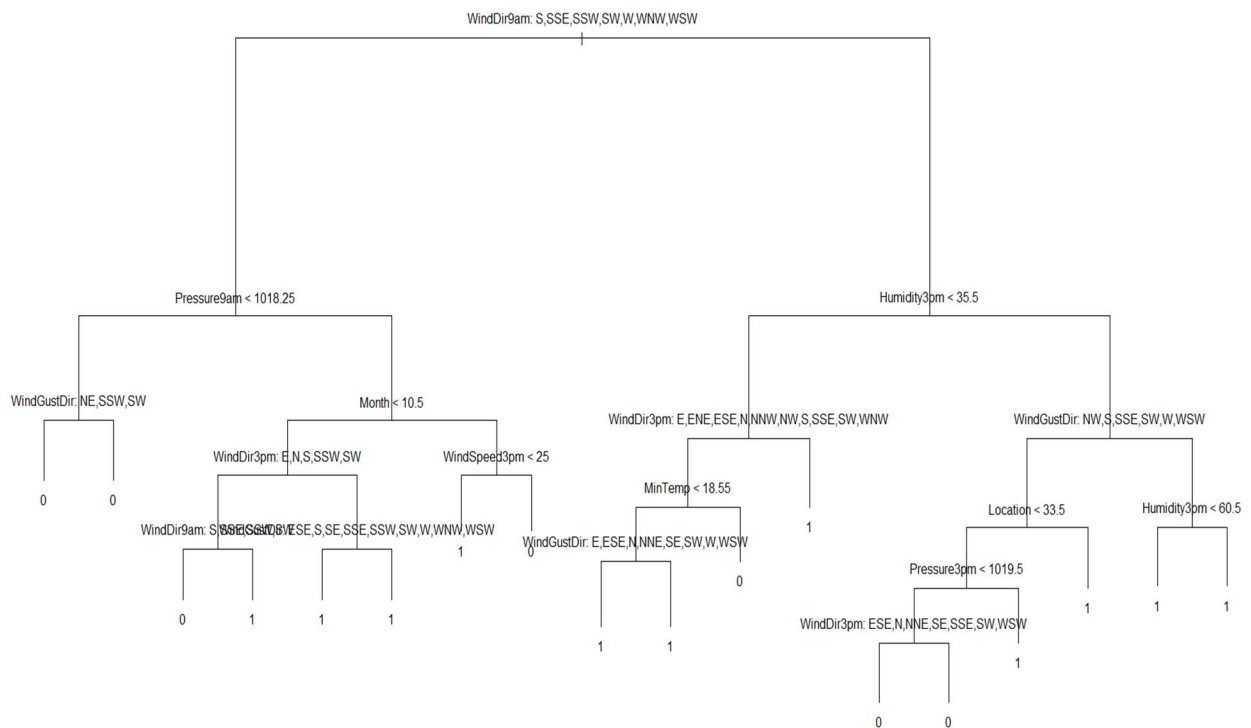
The preprocessing required to make the data set suitable for the model fitting that follows is to firstly remove the 4 attributes with a lot of NA's as mentioned in question 1. After that, since there's still NA's in the remaining attributes, I remove all the rows which contains NA. Screenshot below shows the summary of data after preprocessing and it shows that all remaining attributes has no NA at all.

```
r<-df<-na.omit(df)
r<-summary(df)
```

Day	Month	Year	Location	MinTemp	MaxTemp	Rainfall	windGustDir	windGustSpeed
Min. : 1.0	Min. : 1.00	Min. : 2008	Min. : 7.0	Min. : -2.4	Min. : 9.4	Min. : 0.0	W : 144	Min. : 13.0
1st Qu.: 8.0	1st Qu.: 4.00	1st Qu.: 2011	1st Qu.: 19.0	1st Qu.: 7.3	1st Qu.: 16.8	1st Qu.: 0.0	SSE : 140	1st Qu.: 31.0
Median : 16.0	Median : 7.00	Median : 2014	Median : 23.0	Median : 11.1	Median : 21.7	Median : 0.0	N : 119	Median : 41.0
Mean : 15.6	Mean : 6.86	Mean : 2014	Mean : 22.8	Mean : 11.4	Mean : 22.3	Mean : 1.9	S : 112	Mean : 42.1
3rd Qu.: 23.0	3rd Qu.: 10.00	3rd Qu.: 2017	3rd Qu.: 27.0	3rd Qu.: 15.3	3rd Qu.: 26.7	3rd Qu.: 0.8	WSW : 111	3rd Qu.: 50.0
Max. : 31.0	Max. : 12.00	Max. : 2019	Max. : 34.0	Max. : 26.6	Max. : 44.1	Max. : 128.4	SW : 109	Max. : 111.0
							(other): 671	
windDir9am	windDir3pm	windSpeed9am	windSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm	Temp9am
N : 164	W : 139	Min. : 2.0	Min. : 2	Min. : 10.0	Min. : 4.0	Min. : 994	Min. : 994	Min. : 0.9
SSE : 127	S : 134	1st Qu.: 11.0	1st Qu.: 15	1st Qu.: 57.0	1st Qu.: 37.0	1st Qu.: 1013	1st Qu.: 1011	1st Qu.: 11.7
S : 119	SW : 112	Median : 15.0	Median : 20	Median : 70.0	Median : 54.0	Median : 1018	Median : 1016	Median : 15.5
NNW : 100	WSW : 111	Mean : 16.6	Mean : 21	Mean : 68.9	Mean : 52.4	Mean : 1018	Mean : 1016	Mean : 16.0
NNE : 96	N : 108	3rd Qu.: 20.0	3rd Qu.: 26	3rd Qu.: 83.0	3rd Qu.: 67.0	3rd Qu.: 1023	3rd Qu.: 1021	3rd Qu.: 19.8
SE : 90	SSE : 98	Max. : 63.0	Max. : 61	Max. : 100.0	Max. : 100.0	Max. : 1037	Max. : 1034	Max. : 35.3
(Other): 710	(Other): 704							
Temp3pm	warmerTomorrow							
Min. : 7.9	0: 634							
1st Qu.: 15.6	1: 772							
Median : 20.1								
Mean : 20.7								
3rd Qu.: 24.8								
Max. : 43.0								

### **Question 3 & Question 4:**

For these 2 questions, there is nothing to write for the report, refer to the codes in appendix for more details. The image below is the decision tree created that is being visualized.



### Question 5:

```
#Decision Tree Confusion
r<=print(t1)
      Actual_Class
Predicted_Class  0    1
0      89    72
1     98   163
r<=(89+163)/(89+72+98+163) #accuracy calculation
[1] 0.597
```

Based on the above screenshot which consist of the confusion matrix and accuracy calculation, we can see that the accuracy of decision tree is 0.597.

```
#NaiveBayes Confusion
r<=print(t2)
      Actual_Class
Predicted_Class  0    1
0      99    58
1     88   177
r<=(99+177)/(99+58+88+177)
[1] 0.654
```

Based on the above screenshot, we can see that the accuracy for Naïve Bayes is 0.654.

```
r<=dfpred.bag <- predict.bagging(df.bag, df.test)
r<=dfpred.bag$confusion
      observed class
Predicted Class  0    1
0      102    53
1       85   182
r<=(102+182)/(102+53+85+182)
[1] 0.673
```

Based on the above screenshot, we can see that the accuracy for bagging is 0.673.

```

r<=dfpred.boost <- predict.boosting(df.boost, newdata=df.test)
r<=dfpred.boost$confusion
      observed Class
Predicted Class  0    1
               0 111  60
               1  76 175
r<=(111+175)/(111+60+76+175)
[1] 0.678

```

Based on the above screenshot, we can see that the accuracy for boosting is 0.678.

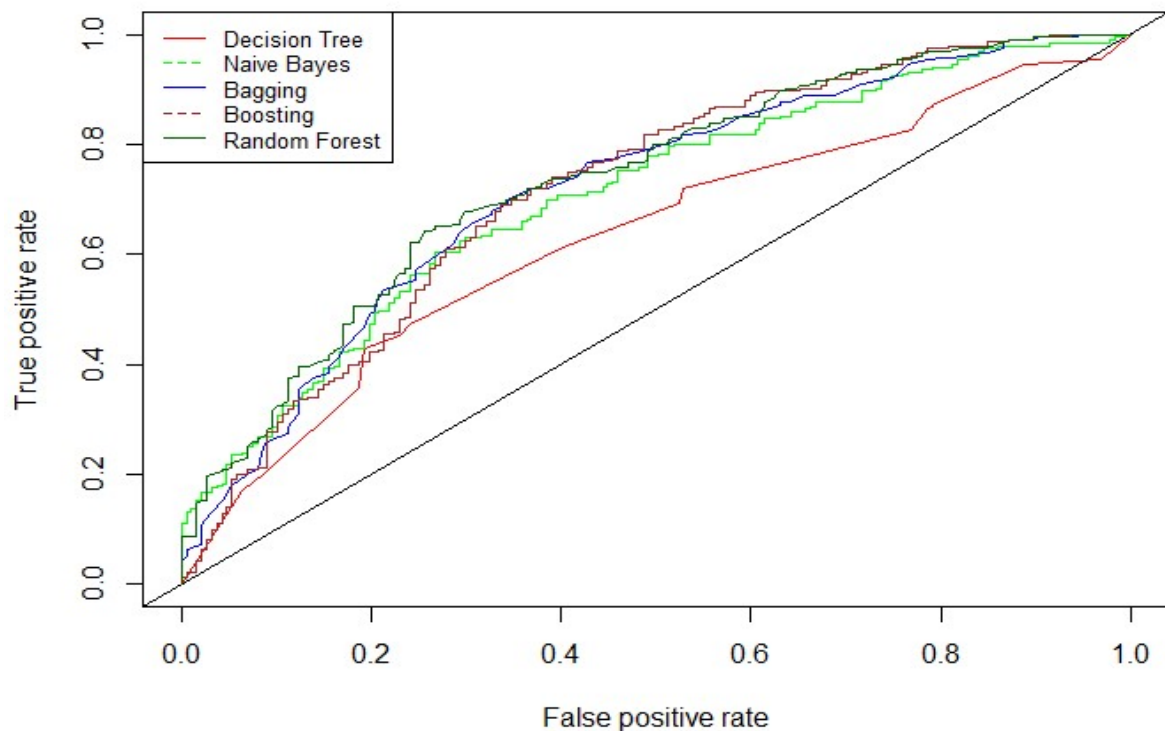
```

#Random Forest Confusion
r<=print(t3)
      Actual_Class
Predicted_Class  0    1
               0  85  40
               1 102 195
r<=(85+195)/(85+40+102+195)
[1] 0.664

```

Based on the above screenshot, we can see that the accuracy for random forest is 0.664.

### Question 6:



The AUC calculated for each classifier is 0.631, 0.706, 0.718, 0.717 and 0.734 respectively for decision tree, Naïve Bayes, Bagging, Boosting and Random Forest respectively.

### Question 7:

```
r<=df2
```

	Classification Method	Area Under Curve	Accuracy
1	Decision Tree	0.631	0.597
2	Naive Bayes	0.706	0.654
3	Bagging	0.718	0.673
4	Boosting	0.717	0.678
5	Random Forest	0.734	0.664

The single best classifier is boosting as it has the higher accuracy and the second highest AUC.

### Question 8:

```
#Decision Tree Attribute Importance
r<=print(summary(df.fit))

Classification tree:
tree(formula = WarmerTomorrow ~ ., data = df.train)
Variables actually used in tree construction:
[1] "windDir9am" "Pressure9am" "windGustDir" "Month" "windDir3pm" "windSpeed3pm" "Humidity3pm" "MinTemp" "Location"
[10] "Pressure3pm"
Number of terminal nodes: 18
Residual mean deviance: 1.02 = 989 / 966
Misclassification error rate: 0.263 = 259 / 984
r<=plot(df.fit)
r<=text(df.fit,pretty=0)
r<=cat("\n#Bagging Attribute Importance\n")

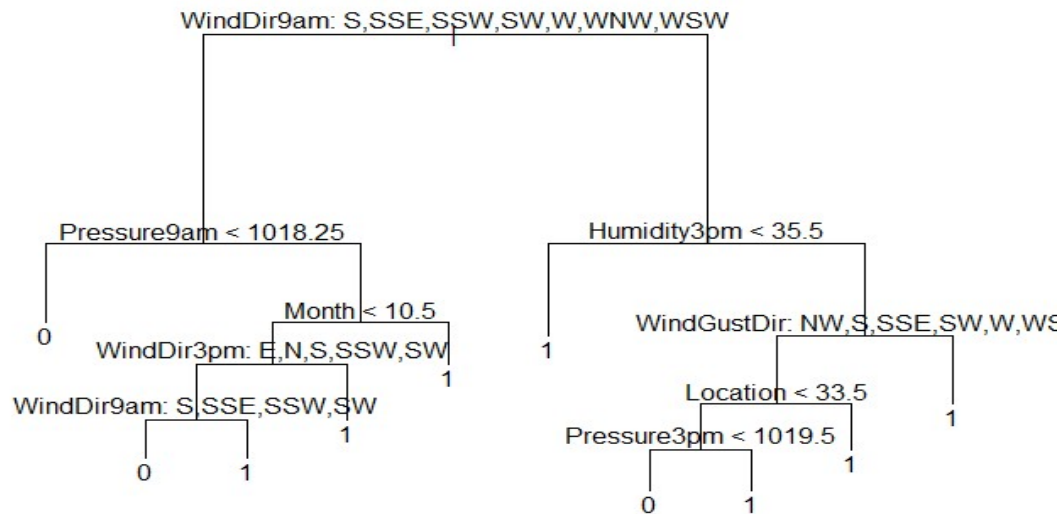
#Bagging Attribute Importance
r<=print(df.bag$importance)
      Day Humidity3pm Humidity9am Location MaxTemp MinTemp Month Pressure3pm Pressure9am Rainfall
1 2.028      6.334      2.119      1.800      4.870      5.912      2.022      1.976      9.352      1.510
2  Temp3pm  Temp9am  windDir3pm  windDir9am  windGustDir  windGustSpeed  windSpeed3pm  windSpeed9am  Year
3 2.812      1.723      16.170      19.660      16.289      1.759      1.182      0.866      1.615
r<=cat("\n#Boosting Attribute Importance\n")

#Boosting Attribute Importance
r<=print(df.boost$importance)
      Day Humidity3pm Humidity9am Location MaxTemp MinTemp Month Pressure3pm Pressure9am Rainfall
1 3.62      3.90      4.81      2.11      6.29      6.02      2.64      2.61      4.64      1.80
2  Temp3pm  Temp9am  windDir3pm  windDir9am  windGustDir  windGustSpeed  windSpeed3pm  windSpeed9am  Year
3 2.60      3.30      14.49      16.11      13.34      2.90      2.81      2.95      3.07
r<=cat("\n#Random Forest Attribute Importance\n")

#Random Forest Attribute Importance
r<=print(df.rf$importance)
      MeanDecreaseGini
Day      17.8
Month    13.3
Year     13.3
Location 12.7
MinTemp  30.6
MaxTemp  27.7
Rainfall 13.4
windGustDir 51.4
windGustSpeed 18.6
windDir9am 56.4
windDir3pm 52.9
windSpeed9am 14.6
windSpeed3pm 17.4
Humidity9am 20.5
Humidity3pm 26.5
Pressure9am 29.9
Pressure3pm 21.4
Temp9am    23.1
Temp3pm    25.7
```

For decision tree, the most important attribute is the WindDir9am as it is the root of the tree, the attributes that can be omitted are those that are not used for tree construction. We can see attributes used for construction of tree through the summary. For bagging, the most important attribute is the WindDir9am as it has the highest value of 19.660. The attribute that can be omitted with least effect on performance will be WindSpeed9am with value of 0.866. For boosting, the most important attribute is the WindDir9am as it has the highest value of 16.11. The attribute that can be omitted with least effect on performance will be Rainfall with value of 1.80. For random forest, the most important attribute is the WindDir9am as it has the highest value of 56.4. The attribute that can be omitted with least effect on performance will be Location with value of 12.7.

**Question 9:**



```

r<=print(t4)
      actual
predicted 0 1
0 87 65
1 100 170
r<=print(t1)
      Actual_Class
Predicted_Class 0 1
0 89 72
1 98 163
r<=#accuracy for t4
r<=(87+170)/(87+65+100+170)
[1] 0.609
r<=#accuracy for t1
r<=(89+163)/(89+72+98+163)
[1] 0.597

```

This model of decision tree performs better than the one in question 4 as it has higher accuracy which is 0.609. However, it is still not performing better than the other 4 classification model in question 4. The important factor in my decision is the relationship between the attributes and WarmerTomorrow. The attributes I used are chosen as they have the highest value of importance level as compare to other attributes.

#### **Question 10:**

```
r<-df.boost2<-boosting(warmerTomorrow~.,df.train,mfinal=500,coflearn = 'Freund')
r<-dfpred.boost2 <- predict.boosting(df.boost2, newdata=df.test)
r<-dfpred.boost2$confusion
      Observed Class
Predicted Class  0    1
      0 111    57
      1  76   178
r<=(111+178)/(111+57+76+178)
[1] 0.685
```

---

As shown above, the accuracy of this boosting model is higher than those in question 4. I chose this model because boosting method was the best in question 4 and hence I would like to improve it. Basically I created this improved model by increasing the mfinal to 500, the default was 100 and I changed the coflearn to Freund so that it uses different weight updating coefficient method. I chose to use all attributes as I tried to use only a few highest important attributes to do it but it had even lower accuracy, so I stick to use all attributes after preprocessing the data.

#### **Question 11:**

The attributes I used are WindDir9am , Humidity3pm , Pressure9am and WindGustDir as these 4 attributes are the more important attributes based on my improved decision tree in question 9. As for the preprocessing required to implement ANN, basically I set the attributes with string values into numeric with values from 0 to 15. The WarmerTomorrow is also turned into numeric as well because it was set into binary.

```
r<=#confusion matrix
r<=table(observed = df.test$warmerTomorrow, predicted = df.pred < 0.5)
      predicted
observed FALSE TRUE
      1   105    82
      2    65   170
r<=(105+170)/(105+82+65+170)
[1] 0.652
```

Based on the screenshot above, the accuracy is 0.652. It is only better than decision tree in terms of performance but it's performance is almost similar to the remaining 4 classification models in question 4. It is probably because the dataset is too small for the neural network to perform better, generally the ANN should perform better than all the 5 classification models in question 4.

**Appendix:**

#Jin En Tan

#Student ID:31336574

#FIT3152 assignment 2

getwd()

setwd("c:/Users/Vapor-15 Pro/Downloads")

#creating individual data set

rm(list = ls())

options(digits = 3,prompt ="r<=")

WAUS <- read.csv("WarmerTomorrow2022.csv",stringsAsFactors = TRUE)

WAUS\$WarmerTomorrow=factor(WAUS\$WarmerTomorrow)

L <- as.data.frame(c(1:49))

set.seed(31336574) # Your Student ID is the random seed

L <- L[sample(nrow(L), 10, replace = FALSE),] # sample 10 locations

WAUS <- WAUS[(WAUS\$Location %in% L),]

WAUS <- WAUS[sample(nrow(WAUS), 2000, replace = FALSE),] # sample 2000 rows

#libraries

library(dplyr)

library(tree)

library(e1071)

library(adabag)

library(randomForest)



#q1

```
WAUS %>% count(WarmerTomorrow) #ignoring those NA
```

```
887/(887+1094)*100 #this is percentage of cooler than previous day
```

```
1094/(887+1094)*100 #this is percentage of warmer than previous day
```

#Based on this we know that the proportion of day which is warmer than previous day is higher than that of

#cooler than previous day.

```
summary(WAUS)
```

#Based on the descriptions of the predictors, we can see that evaporation, sunshine, cloud9am and cloud3pm contains

#a lot of NA. Hence, I consider to remove these 4 attributes from my analysis as there's too many NA in the data.

#q2

#preprocessing, remove those 4 attributes with a lot of NA

```
df<-WAUS[,c(1,2,3,4,5,6,7,10,11,12,13,14,15,16,17,18,19,22,23,24)]
```

#remove remaining rows with NA

```
df<-na.omit(df)
```

```
summary(df)
```

#q3

```
set.seed(31336574) #Student ID as random seed
```

```
train.row = sample(1:nrow(df), 0.7*nrow(df))
```

```
df.train = df[train.row,]
```

```
df.test = df[-train.row,]
```

#q4

#decision tree

```
df.fit=tree(WarmerTomorrow~.,data=df.train)
```

```
plot(df.fit)
text(df.fit,pretty=0)
```

```
#naive bayes
df.naive <- naiveBayes(WarmerTomorrow~, df.train)
```

```
#bagging
df.bag <- bagging(WarmerTomorrow~, df.train)
```

```
#boosting
df.boost<-boosting(WarmerTomorrow~,df.train)
```

```
#random forest
df.rf<-randomForest(WarmerTomorrow~,df.train)
```

```
#q5
#decision tree
df.predtree = predict(df.fit, df.test, type = "class")
t1=table(Predicted_Class = df.predtree, Actual_Class = df.test$WarmerTomorrow)
cat("\n#Decision Tree Confusion\n")
print(t1)
(89+163)/(89+72+98+163) #accuracy calculation
```

```
#naive bayes
df.predbayes = predict(df.naive, df.test)
t2=table(Predicted_Class = df.predbayes, Actual_Class = df.test$WarmerTomorrow)
cat("\n#NaiveBayes Confusion\n")
print(t2)
(99+177)/(99+58+88+177)
```

```
#bagging
```

```
dfpred.bag <- predict.bagging(df.bag, df.test)
```

```
dfpred.bag$confusion
```

```
(102+182)/(102+53+85+182)
```

```
#boosting
```

```
dfpred.boost <- predict.boosting(df.boost, newdata=df.test)
```

```
dfpred.boost$confusion
```

```
(111+175)/(111+60+76+175)
```

```
#random forest
```

```
dfpredrf <- predict(df.rf, df.test)
```

```
t3=table(Predicted_Class = dfpredrf, Actual_Class = df.test$WarmerTomorrow)
```

```
cat("\n#Random Forest Confusion\n")
```

```
print(t3)
```

```
(85+195)/(85+40+102+195)
```

```
#q6
```

```
#decision tree
```

```
# do predictions as probabilities and draw ROC
```

```
library(ROCR)
```

```
df.pred.tree = predict(df.fit, df.test, type = "vector")
```

```
# computing a simple ROC curve (x-axis: fpr, y-axis: tpr)
```

```
# labels are actual values, predictors are probability of class
```

```
dfDpred <- ROCR::prediction(df.pred.tree[,2], df.test$WarmerTomorrow)
```

```
dfDperf <- performance(dfDpred, "tpr", "fpr")
```

```
treeauc = performance(dfDpred, "auc")
```

```
print(as.numeric(treeauc@y.values))
```

```
plot(dfDperf,col="red")
```

```
abline(0,1)
```

```
#naive bayes
```

```
#output as confidence level
```

```
dfpred.bayes = predict(df.naive, df.test, type = 'raw')
```

```
dfBpred <- ROCR::prediction( dfpred.bayes[,2], df.test$WarmerTomorrow)
```

```
dfBperf <- performance(dfBpred,"tpr","fpr")
```

```
naiveauc = performance(dfBpred, "auc")
```

```
print(as.numeric(naiveauc@y.values))
```

```
plot(dfBperf, add=TRUE, col = "green")
```

```
#bagging
```

```
dfBagpred <- ROCR::prediction( dfpred.bag$prob[,2], df.test$WarmerTomorrow)
```

```
dfBagperf <- performance(dfBagpred,"tpr","fpr")
```

```
bagauc = performance(dfBagpred, "auc")
```

```
print(as.numeric(bagauc@y.values))
```

```
plot(dfBagperf, add=TRUE, col = "blue")
```

```
#boosting
```

```
dfBoostpred <- ROCR::prediction( dfpred.boost$prob[,2], df.test$WarmerTomorrow)
```

```
dfBoostperf <- performance(dfBoostpred,"tpr","fpr")
```

```
boostauc = performance(dfBoostpred, "auc")
```

```
print(as.numeric(boostauc@y.values))
```

```
plot(dfBoostperf, add=TRUE, col = "brown")
```

```
#random forest
```

```
dfpred.rf <- predict(df.rf, df.test, type="prob")
```

```
dfFpred <- ROCR::prediction( dfpred.rf[,2], df.test$WarmerTomorrow)
```

```

dfFperf <- performance(dfFpred,"tpr","fpr")

rfauc = performance(dfFpred, "auc")

print(as.numeric(rfauc@y.values))

plot(dfFperf, add=TRUE, col = "darkgreen")

legend("topleft",legend=c("Decision Tree","Naive Bayes","Bagging","Boosting","Random
Forest"),col=c("red","green","blue","brown","darkgreen"), lty=1:2, cex=0.8)

```

#q7

```

comp_auc=c(0.631,0.706,0.718,0.717,0.734)

comp_acc=c(0.597,0.654,0.673,0.678,0.664)

comp_class=c("Decision Tree","Naive Bayes","Bagging","Boosting","Random Forest")

df2<-data.frame(comp_class,comp_auc,comp_acc)

colnames(df2)<-c("Classification Method","Area Under Curve","Accuracy")

df2

#the single best classifier is boosting as it has the highest accuracy and a relatively higher area under
curve as well

```

#q8

```

#Attribute importance

cat("\n#Decision Tree Attribute Importance\n")

print(summary(df.fit))

plot(df.fit)

text(df.fit,pretty=0)

cat("\n#Bagging Attribute Importance\n")

print(df.bag$importance)

cat("\n#Boosting Attribute Importance\n")

print(df.boost$importance)

cat("\n#Random Forest Attribute Importance\n")

print(df.rf$importance)

```

#for decision tree, the most important attribute is the WindDir9am as it is the root of the tree,  
#the attributes that can be omitted are those that are not used for tree construction. We can see attributes used for  
#construction of tree through the summary  
#for bagging, the most important attribute is the WindDir9am as it has the highest value of 19.660  
#the attribute that can be omitted with least effect on performance will be WindSpeed9am with value of 0.866  
#for boosting, the most important attribute is the WindDir9am as it has the highest value of 16.11  
#the attribute that can be omitted with least effect on performance will be Rainfall with value of 1.80  
#for random forest, the most important attribute is the WindDir9am as it has the highest value of 56.4  
#the attribute that can be omitted with least effect on performance will be Location with value of 12.7

#q9

#cross validation and pruning

```
test.fit=cv.tree(df.fit, FUN=prune.misclass)
print(test.fit)
prune.dffit = prune.misclass(df.fit, best=10)
print(summary(prune.dffit))
plot(prune.dffit)
text(prune.dffit, pretty=0)
#test accuracy after pruning
dfp.predict = predict(prune.dffit, df.test, type = "class")
t4=table(predicted = dfp.predict, actual = df.test$WarmerTomorrow)
print(t4)
print(t1)
#accuracy for t4
(87+170)/(87+65+100+170)
#accuracy for t1
(89+163)/(89+72+98+163)
```

#this model is better than the previous model in part 4 as the accuracy is higher

#the relationship between attributes and the WarmerTomorrow is important in the decision

#these attributes are chosen because the importance level is among the highest as compare to other attributes

#q10

#boosting

```
df.boost2<-boosting(WarmerTomorrow~,df.train,mfinal=500,coflearn = 'Freund')
```

```
dfpred.boost2 <- predict.boosting(df.boost2, newdata=df.test)
```

```
dfpred.boost2$confusion
```

```
(111+178)/(111+57+76+178)
```

```
library(neuralnet)
```

#q11

#preprocessing

```
df.train$WindGustDir=recode(df.train$WindGustDir,'E'='0','ENE'='1','ESE'='2','N'='3','NE'='4','NNE'='5','NNW'='6','NW'='7','S'='8','SE'='9','SSE'='10','SSW'='11','SW'='12','W'='13','WNW'='14','WSW'='15')
```

```
df.test$WindGustDir=recode(df.test$WindGustDir,'E'='0','ENE'='1','ESE'='2','N'='3','NE'='4','NNE'='5','NNW'='6','NW'='7','S'='8','SE'='9','SSE'='10','SSW'='11','SW'='12','W'='13','WNW'='14','WSW'='15')
```

```
df.train$WindDir9am=recode(df.train$WindDir9am,'E'='0','ENE'='1','ESE'='2','N'='3','NE'='4','NNE'='5','NNW'='6','NW'='7','S'='8','SE'='9','SSE'='10','SSW'='11','SW'='12','W'='13','WNW'='14','WSW'='15')
```

```
df.test$WindDir9am=recode(df.test$WindDir9am,'E'='0','ENE'='1','ESE'='2','N'='3','NE'='4','NNE'='5','NNW'='6','NW'='7','S'='8','SE'='9','SSE'='10','SSW'='11','SW'='12','W'='13','WNW'='14','WSW'='15')
```

```
df.train$WindDir3pm=recode(df.train$WindDir3pm,'E'='0','ENE'='1','ESE'='2','N'='3','NE'='4','NNE'='5','NNW'='6','NW'='7','S'='8','SE'='9','SSE'='10','SSW'='11','SW'='12','W'='13','WNW'='14','WSW'='15')
```

```
df.test$WindDir3pm=recode(df.test$WindDir3pm,'E'='0','ENE'='1','ESE'='2','N'='3','NE'='4','NNE'='5','NNW'='6','NW'='7','S'='8','SE'='9','SSE'='10','SSW'='11','SW'='12','W'='13','WNW'='14','WSW'='15')
```

```
df.train$WarmerTomorrow <- as.numeric(df.train$WarmerTomorrow)
```

```
df.train$WindGustDir <- as.numeric(df.train$WindGustDir)
```

```

df.train$WindDir9am <- as.numeric(df.train$WindDir9am)
df.train$WindDir3pm <- as.numeric(df.train$WindDir3pm)
df.test$WarmerTomorrow <- as.numeric(df.test$WarmerTomorrow)
df.test$WindGustDir <- as.numeric(df.test$WindGustDir)
df.test$WindDir9am <- as.numeric(df.test$WindDir9am)
df.test$WindDir3pm <- as.numeric(df.test$WindDir3pm)

df.nn = neuralnet(WarmerTomorrow == 1 ~ WindDir9am + Humidity3pm + Pressure9am + WindGustDir,
df.train,
                hidden=1,linear.output = FALSE)
df.pred = predict(df.nn, df.test)
#confusion matrix
table(observed = df.test$WarmerTomorrow, predicted = df.pred < 0.5)
(105+170)/(105+82+65+170)

```